

# Non-linear Models for Classification

04/30/2021

Seungheon Han

1. Use KNN & SVM to predict whether a given car gets high or low gas mileage based on the Auto data set.

- (a) Create a binary variable mpg01, and generate the training and test sets

```
pacman::p_load(ISLR, GGally, ggplot2, dplyr, class, e1071, caret, tree, gbm, randomForest)
```

```
auto <- Auto[, setdiff(names(Auto), c("name"))]
auto$origin <- as.factor(auto$origin)
auto$cylinders <- as.factor(auto$cylinders)
str(auto)
```

```
## 'data.frame':   392 obs. of  8 variables:
## $ mpg          : num  18 15 18 16 17 15 14 14 15 ...
## $ cylinders    : Factor w/ 5 levels "3","4","5","6",...: 5 5 5 5 5 5 5 5 5 ...
## $ displacement: num  307 350 318 304 302 429 454 440 455 390 ...
## $ horsepower  : num  130 165 150 150 140 198 220 215 225 190 ...
## $ weight       : num  3504 3693 3436 3433 3449 ...
## $ acceleration: num  12 11.5 11 12 10.5 10 9 8.5 10 8.5 ...
## $ year        : num  70 70 70 70 70 70 70 70 70 70 ...
## $ origin      : Factor w/ 3 levels "1","2","3": 1 1 1 1 1 1 1 1 1 ...
```

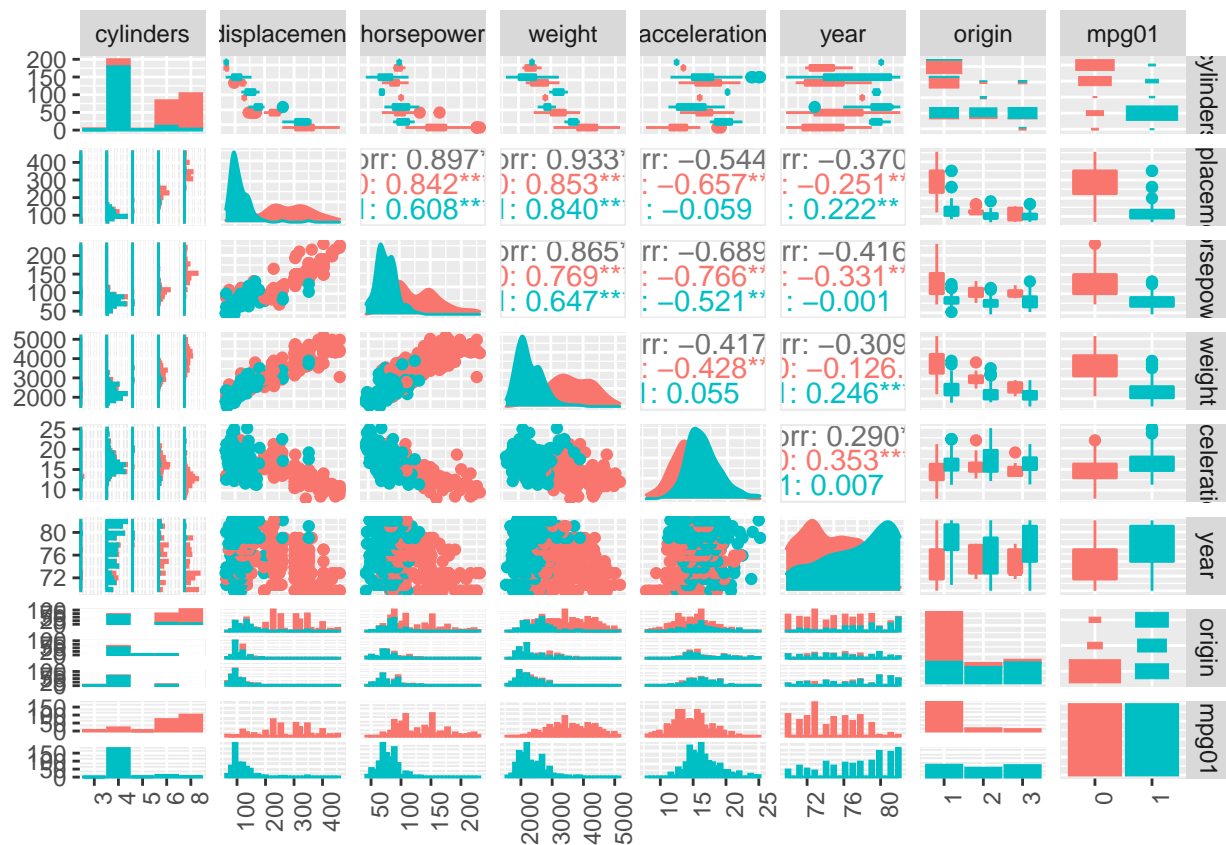
```
auto$mpg01 <- ifelse(auto$mpg > median(auto$mpg), 1, 0)
auto <- auto[, setdiff(names(auto), c("mpg"))]
auto$mpg01 <- as.factor(auto$mpg01)
```

```
set.seed(123)
split <- sample(2, nrow(auto), replace=T, prob = c(0.8, 0.2))
train <- auto[split == 1,]
test <- auto[split == 2,]
```



- (b) Perform KNN on the training data, with several values of  $K$ , in order to predict mpg01. Use only the variables that seemed most associated with mpg01.



```
### exploring the relationships between mpg01 and each of the predictor
auto %>%
  ggpairs(aes(col = mpg01, fill = mpg01)) +
  theme(axis.text.x = element_text(angle = 90, hjust = 1))
```



```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```





```
for(i in 1:ncol(auto)){
  print(ggplot(auto, aes_string("mpg01", colnames(auto)[i]))+
    geom_jitter(aes(col=origin)))
}
```









```

### the most associated predictors: cylinders, displacement, horsepower, weight, year
train <- train %>%
  select(-c("acceleration", "origin"))
test <- test %>%
  select(-c("acceleration", "origin"))

train_x <- as.data.frame(train[,setdiff(names(train), c("mpg01"))])
train_y <- train[, "mpg01"]
test_x <- as.data.frame(test[, setdiff(names(test), c("mpg01"))])
test_y <- test[, "mpg01"]
### k-fold CV for find the best K in KNN model
set.seed(123)
K.knn = c(1:20)
k.fold <- 10
folds <- createFolds(train_y, k = k.fold)

error.vec = NULL
for(i in K.knn){
  error.list = NULL
  for(j in 1:k.fold){
    knn.pred <- knn(train_x[-unlist(folds[j]),], train_x[unlist(folds[j]),], train_y[-unlist(folds[j])])
    cv.error <- sum(knn.pred != train_y[unlist(folds[j])])/length(unlist(folds[j]))
    error.list <- c(error.list, cv.error)
  }
  avg.error <- round(mean(error.list),4)
  error.vec <- c(error.vec, avg.error)
  print(paste("K =", i, "cv-error:", avg.error))
}

```

```

## [1] "K = 1 cv-error: 0.1226"
## [1] "K = 2 cv-error: 0.1354"
## [1] "K = 3 cv-error: 0.1356"
## [1] "K = 4 cv-error: 0.1261"
## [1] "K = 5 cv-error: 0.1163"
## [1] "K = 6 cv-error: 0.1226"
## [1] "K = 7 cv-error: 0.1226"
## [1] "K = 8 cv-error: 0.1162"
## [1] "K = 9 cv-error: 0.1228"
## [1] "K = 10 cv-error: 0.126"
## [1] "K = 11 cv-error: 0.1259"
## [1] "K = 12 cv-error: 0.1386"
## [1] "K = 13 cv-error: 0.1355"
## [1] "K = 14 cv-error: 0.1355"
## [1] "K = 15 cv-error: 0.1386"
## [1] "K = 16 cv-error: 0.1385"
## [1] "K = 17 cv-error: 0.1449"
## [1] "K = 18 cv-error: 0.1353"
## [1] "K = 19 cv-error: 0.1353"
## [1] "K = 20 cv-error: 0.1385"

```

```

print(paste("The K minimizing the CV error: K =", match(min(error.vec), error.vec), "with error", min(e

```

```

## [1] "The K minimizing the CV error: K = 8 with error 0.1162"

```

```
# tune.knn function generates the same result
set.seed(123)
tune.out.knn <- tune.knn(x = train_x, y = train_y, k=1:20)
print(paste("Best K:", tune.out.knn$best.parameters))
```

```
## [1] "Best K: 8"
```

```
### comment ###
# K=8 generates the smallest validation error.
```

- (c) Fit a support vector classifier to the training data with various values of `cost`, in order to predict whether a car gets high or low gas mileage.

```
### (c) ###
set.seed(123)
tune.out.lin <- tune(svm, mpg01~., data = train, kernal = "linear",
                    ranges = list(cost = c(0.1, 1, 10, 100, 1000, 10000, 30000)))

### CV-errors by different costs
summary(tune.out.lin)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   10000
##
## - best performance: 0.08165323
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-01 0.10352823 0.05078718
## 2 1e+00 0.10665323 0.05516446
## 3 1e+01 0.08800403 0.04363379
## 4 1e+02 0.08487903 0.02960729
## 5 1e+03 0.08165323 0.02604546
## 6 1e+04 0.08165323 0.03015969
## 7 3e+04 0.09465726 0.05679151
```

```
### best cost
bestmod <- tune.out.lin$best.model
print(paste("the best cost is", summary(bestmod)$cost))
```

```
## [1] "the best cost is 10000"
```

```
### fit the svm model
svm.fit.lin <- svm(formula= mpg01~., data = train, kernal = "linear",
                  cost = 10000)
summary(svm.fit.lin)
```



```
##
## Call:
## svm(formula = mpg01 ~ ., data = train, kernal = "linear", cost = 10000)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##         cost: 10000
##
## Number of Support Vectors:  51
##
## ( 29 22 )
##
##
## Number of Classes:  2
##
## Levels:
##  0 1
```

Comments:  
with support vector classifier, the best cost is 10000

- (d) Now repeat (c), this time using SVMs with radial and polynomial basis kernels, with different values of gamma and degree and cost.

```
### kernal = radial
set.seed(123)
tune.out.rad <- tune(svm, mpg01~., data=train, kernel="radial",
                    ranges=list(cost=c(0.1,1,10,100,1000),
                                gamma=c(0.5,1,2,3,4) ))

summary(tune.out.rad)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##     1     3
##
## - best performance: 0.06270161
##
## - Detailed performance results:
##   cost gamma      error dispersion
## 1 1e-01   0.5 0.10040323 0.07010733
## 2 1e+00   0.5 0.09102823 0.04960171
## 3 1e+01   0.5 0.07852823 0.02179456
## 4 1e+02   0.5 0.07237903 0.02984280
## 5 1e+03   0.5 0.07570565 0.04797571
## 6 1e-01   1.0 0.09415323 0.05285737
```

```
## 7 1e+00 1.0 0.07842742 0.04461323
## 8 1e+01 1.0 0.06905242 0.03532623
## 9 1e+02 1.0 0.06633065 0.04146783
## 10 1e+03 1.0 0.08820565 0.04911492
## 11 1e-01 2.0 0.09727823 0.04730450
## 12 1e+00 2.0 0.06592742 0.03729748
## 13 1e+01 2.0 0.07883065 0.03455569
## 14 1e+02 2.0 0.08205645 0.04071882
## 15 1e+03 2.0 0.09758065 0.04587598
## 16 1e-01 3.0 0.10362903 0.04645077
## 17 1e+00 3.0 0.06270161 0.03879790
## 18 1e+01 3.0 0.06945565 0.05156250
## 19 1e+02 3.0 0.07883065 0.05236073
## 20 1e+03 3.0 0.09133065 0.05048202
## 21 1e-01 4.0 0.13235887 0.05948065
## 22 1e+00 4.0 0.06602823 0.04996957
## 23 1e+01 4.0 0.06945565 0.05561222
## 24 1e+02 4.0 0.08195565 0.05424859
## 25 1e+03 4.0 0.08820565 0.04685362
```

```
### best gamma & cost
bestmod2 <- tune.out.rad$best.parameters
print(paste("best cost:", bestmod2$cost, "& best gamma:", bestmod2$gamma))
```

```
## [1] "best cost: 1 & best gamma: 3"
```

```
### kernal = polynomial
set.seed(123)
tune.out.poly <- tune(svm, mpg01~., data = train, kernal = "polynomial",
                      ranges = list(cost = c(0.1,1,10,100,1000),
                                     degree = c(1:10)))
summary(tune.out.poly)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost degree
## 1000      1
##
## - best performance: 0.08165323
##
## - Detailed performance results:
##   cost degree      error dispersion
## 1 1e-01      1 0.10352823 0.05078718
## 2 1e+00      1 0.10665323 0.05516446
## 3 1e+01      1 0.08800403 0.04363379
## 4 1e+02      1 0.08487903 0.02960729
## 5 1e+03      1 0.08165323 0.02604546
## 6 1e-01      2 0.10352823 0.05078718
## 7 1e+00      2 0.10665323 0.05516446
```

```
## 8 1e+01      2 0.08800403 0.04363379
## 9 1e+02      2 0.08487903 0.02960729
## 10 1e+03     2 0.08165323 0.02604546
## 11 1e-01     3 0.10352823 0.05078718
## 12 1e+00     3 0.10665323 0.05516446
## 13 1e+01     3 0.08800403 0.04363379
## 14 1e+02     3 0.08487903 0.02960729
## 15 1e+03     3 0.08165323 0.02604546
## 16 1e-01     4 0.10352823 0.05078718
## 17 1e+00     4 0.10665323 0.05516446
## 18 1e+01     4 0.08800403 0.04363379
## 19 1e+02     4 0.08487903 0.02960729
## 20 1e+03     4 0.08165323 0.02604546
## 21 1e-01     5 0.10352823 0.05078718
## 22 1e+00     5 0.10665323 0.05516446
## 23 1e+01     5 0.08800403 0.04363379
## 24 1e+02     5 0.08487903 0.02960729
## 25 1e+03     5 0.08165323 0.02604546
## 26 1e-01     6 0.10352823 0.05078718
## 27 1e+00     6 0.10665323 0.05516446
## 28 1e+01     6 0.08800403 0.04363379
## 29 1e+02     6 0.08487903 0.02960729
## 30 1e+03     6 0.08165323 0.02604546
## 31 1e-01     7 0.10352823 0.05078718
## 32 1e+00     7 0.10665323 0.05516446
## 33 1e+01     7 0.08800403 0.04363379
## 34 1e+02     7 0.08487903 0.02960729
## 35 1e+03     7 0.08165323 0.02604546
## 36 1e-01     8 0.10352823 0.05078718
## 37 1e+00     8 0.10665323 0.05516446
## 38 1e+01     8 0.08800403 0.04363379
## 39 1e+02     8 0.08487903 0.02960729
## 40 1e+03     8 0.08165323 0.02604546
## 41 1e-01     9 0.10352823 0.05078718
## 42 1e+00     9 0.10665323 0.05516446
## 43 1e+01     9 0.08800403 0.04363379
## 44 1e+02     9 0.08487903 0.02960729
## 45 1e+03     9 0.08165323 0.02604546
## 46 1e-01    10 0.10352823 0.05078718
## 47 1e+00    10 0.10665323 0.05516446
## 48 1e+01    10 0.08800403 0.04363379
## 49 1e+02    10 0.08487903 0.02960729
## 50 1e+03    10 0.08165323 0.02604546
```

```
bestmod3 <- tune.out.poly$best.parameters
print(paste("best cost:", bestmod3$cost, "& best degree:", bestmod3$degree))
```


```
## [1] "best cost: 1000 & best degree: 1"
```

Comments:

for SVM with radial kernel, 10-fold CV finds the optimal cost and gamma which have turned out to be 1 and 3 respectively. for SVM with polynomial kernel, the best cost and degree are 1000 and 1 respectively.


(e) Make some plots to back up your assertions in (b), (c) and (d).

```
### best K minimizing the validation error is 8  
plot(tune.out.knn)
```




```
### best cost minimizing the validation error is 10000  
plot(tune.out.lin)
```

### Performance of 'svm'




```
### best cost is 100 and gamma is 3  
plot(tune.out.rad)
```

### Performance of 'svm'



```
### best cost is 10000 and degree is 1
plot(tune.out.poly)
```



(f) Compare the test errors of the best tuned models for KNN, linear SVM, SVM with radial basis kernel, and SVM with polynomial basis kernel.

```
### (f) ###

### KNN model with K=8
optimal.knn <- knn(train_x, test_x, train_y, k=8)
table.knn <- table(optimal.knn, test_y)
error.knn <- 1-sum(diag(table.knn))/sum(table.knn)

### Support Vector Classifier with cost = 10000
svm.fit.lin <- svm(formula= mpg01~., data = train, kernel = "linear",
                  cost = 10000)
pred.lin.svm <- predict(svm.fit.lin, test_x)
table.lin.svm <- table(pred.lin.svm, test_y)
error.lin.svm <- 1-sum(diag(table.lin.svm))/sum(table.lin.svm)

### SVM with radial kernels with cost = 100, gamma = 3
svm.fit.rad <- svm(formula= mpg01~., data = train, kernel = "radial",
                  cost = 100, gamma = 3)
pred.rad.svm <- predict(svm.fit.rad, test_x)
table.rad.svm <- table(pred.rad.svm, test_y)
error.rad.svm <- 1-sum(diag(table.rad.svm))/sum(table.rad.svm)

### SVM with polynomial kernels with cost = 10000, degree = 1
```

```

svm.fit.poly <- svm(formula= mpg01~., data = train, kernal = "polynomial",
                    cost = 1000, degree = 1)
pred.poly.svm <- predict(svm.fit.poly, test_x)
table.poly.svm <- table(pred.poly.svm, test_y)
error.poly.svm <- 1-sum(diag(table.poly.svm))/sum(table.poly.svm)

```

```

### Comparing the test errors generated from the four models
print(paste("test error of knn model:", round(error.knn,4)))

```

```
## [1] "test error of knn model: 0.1486"
```

```
print(paste("test error of linear svm model:", round(error.lin.svm,4)))
```

```
## [1] "test error of linear svm model: 0.0946"
```

```
print(paste("test error of radial svm model:", round(error.rad.svm,4)))
```

```
## [1] "test error of radial svm model: 0.1081"
```

```
print(paste("test error of polynomial svm model:", round(error.poly.svm,4)))
```

```
## [1] "test error of polynomial svm model: 0.0811"
```

```

### comment:
# the SVM model with polynomial basis kernel produces the smallest test error.

```

2. (Regression Tree, Boosting, Bagging and Random Forest) Use the OJ data set which is part of the ISLR package.

(a) Create a training set containing a random sample of 800 observations, and a test set containing the remaining observations.

```

oj <- OJ
set.seed(1)
split.train <- sample(1:nrow(oj), 800)
train <- oj[split.train,]
test <- oj[-split.train,]

```

(b) Fit a tree to the training data, with `Purchase` as the response and the other variables as predictors. Use the `summary()` function to produce summary statistics about the tree, and describe the results obtained.

```

tree.oj.train <- tree(Purchase ~., train)
summary(tree.oj.train)

```

```

##
## Classification tree:
## tree(formula = Purchase ~ ., data = train)

```

```
## Variables actually used in tree construction:
## [1] "LoyalCH"      "PriceDiff"    "SpecialCH"    "ListPriceDiff"
## [5] "PctDiscMM"
## Number of terminal nodes: 9
## Residual mean deviance: 0.7432 = 587.8 / 791
## Misclassification error rate: 0.1588 = 127 / 800
```

Comments:

terminal nodes = 9, error rate = 15.875%

(c) Pick one of the terminal nodes, and interpret the information displayed.

```
tree.oj.train
```

```
## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
## 1) root 800 1073.00 CH ( 0.60625 0.39375 )
##    2) LoyalCH < 0.5036 365 441.60 MM ( 0.29315 0.70685 )
##      4) LoyalCH < 0.280875 177 140.50 MM ( 0.13559 0.86441 )
##        8) LoyalCH < 0.0356415 59 10.14 MM ( 0.01695 0.98305 ) *
##        9) LoyalCH > 0.0356415 118 116.40 MM ( 0.19492 0.80508 ) *
##      5) LoyalCH > 0.280875 188 258.00 MM ( 0.44149 0.55851 )
##        10) PriceDiff < 0.05 79 84.79 MM ( 0.22785 0.77215 )
##          20) SpecialCH < 0.5 64 51.98 MM ( 0.14062 0.85938 ) *
##          21) SpecialCH > 0.5 15 20.19 CH ( 0.60000 0.40000 ) *
##        11) PriceDiff > 0.05 109 147.00 CH ( 0.59633 0.40367 ) *
##    3) LoyalCH > 0.5036 435 337.90 CH ( 0.86897 0.13103 )
##      6) LoyalCH < 0.764572 174 201.00 CH ( 0.73563 0.26437 )
##        12) ListPriceDiff < 0.235 72 99.81 MM ( 0.50000 0.50000 )
##          24) PctDiscMM < 0.196196 55 73.14 CH ( 0.61818 0.38182 ) *
##          25) PctDiscMM > 0.196196 17 12.32 MM ( 0.11765 0.88235 ) *
##      13) ListPriceDiff > 0.235 102 65.43 CH ( 0.90196 0.09804 ) *
##      7) LoyalCH > 0.764572 261 91.20 CH ( 0.95785 0.04215 ) *
```


Comments:

node number (8): the observations whose LoyalCH is smaller than 0.0356415 is classified from the upper branch into the node (8). the number of observations fell in the node 8 is 59, and its deviance is 10.14. the overall prediction for the node is "MM" and the fraction of MM here is 0.98305.

(d) Create a plot of the tree, and interpret the results.

```
plot(tree.oj.train)
text(tree.oj.train, pretty=0)
```





Comments:

the predictors used for the tree is LoyalCH, PriceDiff, SpecialCH, ListPriceDiff, PctDiscMM.

among the predictors, LoyalCH located on the top is the most important variable and the tree is consisting of 8 subtrees with 9 terminal nodes

- (e) Predict the response on the test data, and produce a confusion matrix comparing the test labels to the predicted test labels.

```
tree.pred <- predict(tree.oj.train, test, type = "class")
conf.table <- table(tree.pred, test$Purchase)
print(conf.table)
```

```
##
## tree.pred  CH  MM
##           CH 160  38
##           MM   8  64
```

```
test.error <- 1-sum(diag(conf.table))/sum(conf.table)
print(paste("test error rate: ",round(test.error, 4)))
```


```
## [1] "test error rate:  0.1704"
```

- (f) Determine the optimal tree size.

```
set.seed(1)
cv.tree.oj <- cv.tree(tree.oj.train, FUN = prune.misclass)
```

- (g) Produce a plot with tree size on the  $x$ -axis and cross-validated classification error rate on the  $y$ -axis.

```
plot(cv.tree.oj$size, cv.tree.oj$dev, type = "b")
```




Comments:

based on the cross-validation error for each size, we can decide that the best size(# terminal nodes) should be 8 or 9 because these have the smallest CV error which is 145.

- (i) Produce a pruned tree corresponding to the optimal tree size obtained using cross-validation.

```
## using 8 as the optimal size
prune.oj <- prune.misclass(tree.oj.train, best = 8)
plot(prune.oj)
text(prune.oj, pretty = 0)
```



- (j) Compare the training error rates between the pruned and unpruned trees.

```
summary(tree.oj.train)
```

```
##
## Classification tree:
## tree(formula = Purchase ~ ., data = train)
## Variables actually used in tree construction:
## [1] "LoyalCH"      "PriceDiff"    "SpecialCH"    "ListPriceDiff"
## [5] "PctDiscMM"
## Number of terminal nodes: 9
## Residual mean deviance: 0.7432 = 587.8 / 791
## Misclassification error rate: 0.1588 = 127 / 800
```

```
summary(prune.oj)
```

```
##
## Classification tree:
## snip.tree(tree = tree.oj.train, nodes = 4L)
## Variables actually used in tree construction:
## [1] "LoyalCH"      "PriceDiff"    "SpecialCH"    "ListPriceDiff"
## [5] "PctDiscMM"
## Number of terminal nodes: 8
## Residual mean deviance: 0.7598 = 601.8 / 792
## Misclassification error rate: 0.1588 = 127 / 800
```

Comments:

training error rate of the unpruned model is 15.875%  
training error rate of the pruned model is 15.875%  
- the two models have the same training error rate

(k) Compare the test error rates between the pruned and unpruned trees.

```
pred.pruned <- predict(prune.oj, test, type = "class")
conf.table.pruned <- table(pred.pruned, test$Purchase)
print(conf.table.pruned)
```

```
##
## pred.pruned CH MM
##           CH 160 38
##           MM   8 64
```

```
test.error.pruned <- 1-sum(diag(conf.table.pruned))/sum(conf.table.pruned)
print(paste("test error rate of the unpruned model: ", test.error))
```

```
## [1] "test error rate of the unpruned model: 0.17037037037037"
```

```
print(paste("test error rate of the pruned model: ", test.error.pruned))
```

```
## [1] "test error rate of the pruned model: 0.17037037037037"
```

Comments:

the test error rates of the two models are exactly equal


- (l) Perform boosting on the training set with 1,000 trees for a range of values of the shrinkage parameter  $\lambda$ . Produce a plot with different shrinkage values on the  $x$ -axis and the corresponding training error and test error on the  $y$ -axis.

```
train$Purchase <- ifelse(train$Purchase == "CH", 1, 0)
test$Purchase <- ifelse(test$Purchase == "CH", 1, 0)

## Plot with different shrinkage parameters and train errors
set.seed(1)
shrink <- c(0.1, 0.3, 0.5, 0.7, 0.9)
train.errors <- NULL

for(i in shrink){
  boost2 <- gbm(Purchase~., data = train, distribution = "bernoulli",
               n.trees = 1000, interaction.depth = 4, shrinkage=i)
  pred2 <- ifelse(predict(boost2, newdata = train,
                         n.trees=1000, type="response")>0.5,1,0)
  table2 <- table(pred2, train$Purchase)
  train.errors <- c(train.errors, 1-sum(diag(table2))/sum(table2))
}


train.err.df<- data.frame(train.errors, shrink)
ggplot(train.err.df, aes(x=shrink, y=train.errors))+
  geom_line(color = "blue")+
  geom_point(color = "blue")
```



```
## Plot with different shrinkage parameters and test errors
set.seed(1)
test.errors <- NULL

for(i in shrink){
  boost1 <- gbm(Purchase~., data = train, distribution = "bernoulli",
               n.trees = 1000, interaction.depth = 4, shrinkage=i)
  pred1 <- ifelse(predict(boost1, newdata = test,
                        n.trees=1000, type="response")>0.5,1,0)
  table1 <- table(pred1, test$Purchase)
  test.errors <- c(test.errors, 1-sum(diag(table1))/sum(table1))
}


test.err.df<- data.frame(test.errors, shrink)
ggplot(test.err.df, aes(x=shrink, y=test.errors))+
  geom_line(color = "blue")+
  geom_point(color = "blue")
```



*#results: shrinkage parameter 0.3 minimizes test error.*

```
set.seed(1)
boost.oj.train <- gbm(Purchase~., data = train, distribution = "bernoulli",
                      n.trees = 1000, interaction.depth = 4, shrinkage = 0.3)

summary(boost.oj.train)
```



```
##           var    rel.inf
## LoyalCH      LoyalCH 39.0157396
## WeekofPurchase WeekofPurchase 18.9430929
## PriceDiff    PriceDiff 9.0393249
## StoreID      StoreID 5.5850013
## ListPriceDiff ListPriceDiff 4.9479441
## STORE        STORE 4.7484284
## SalePriceMM   SalePriceMM 4.2228905
## PriceCH       PriceCH 2.8743455
## PriceMM       PriceMM 2.6483894
## SalePriceCH   SalePriceCH 2.3895021
## SpecialCH     SpecialCH 1.7332148
## DiscMM        DiscMM 1.3380080
## PctDiscMM     PctDiscMM 0.8448488
## SpecialMM     SpecialMM 0.7576262
## DiscCH        DiscCH 0.3626796
## Store7        Store7 0.2956998
## PctDiscCH     PctDiscCH 0.2532640
```

```
pred.boost <- ifelse(predict(boost.oj.train, newdata = test,
                             n.trees=1000, type="response")>0.5,1,0)
table1 <- table(pred.boost, test$Purchase)
boost.test.error <- 1-sum(diag(table1))/sum(table1)

print(paste("test error of the boosting model with shrinkage parameter = 0.3 is: ", round(boost.test.er
```

```
## [1] "test error of the boosting model with shrinkage parameter = 0.3 is: 0.1963"
```

Comments:

the important predictors of this model are:  
 "LoyalCH" and "WeekofPurchase"

(m) Perform bagging on the training set and report the prediction performance on the test set.

```
set.seed(1)
split.train <- sample(1:nrow(oj), 800)
train <- oj[split.train,]
test <- oj[-split.train,]

set.seed(1)
bag.oj.train = randomForest(Purchase~., data = train, mtry=17, importance=TRUE)
pred.bag <- predict(bag.oj.train, newdata = test)
bag.table <- table(pred.bag, test$Purchase)
bag.error <- 1-sum(diag(bag.table))/sum(bag.table)
print(paste("test error of the bagging model: ", round(bag.error, 4)))
```

```
## [1] "test error of the bagging model: 0.1852"
```

```
bag.oj.train$importance
```

##		CH	MM	MeanDecreaseAccuracy	MeanDecreaseGini
## WeekofPurchase	1.449275e-02	0.0081761108		0.0119616206	38.3886426
## StoreID	2.813469e-03	0.0157406190		0.0078979001	12.6325113
## PriceCH	3.908277e-03	0.0031710145		0.0036595727	4.3656852
## PriceMM	2.065023e-03	0.0032540724		0.0025174091	4.4250917
## DiscCH	-6.707993e-04	0.0024988832		0.0005629286	1.7870508
## DiscMM	1.090828e-03	0.0032210479		0.0019057516	2.1431513
## SpecialCH	2.555862e-03	0.0043654998		0.0033002920	6.5637446
## SpecialMM	-5.084022e-05	0.0010224016		0.0003630224	2.6593366
## LoyalCH	1.370922e-01	0.2559065843		0.1835532385	227.1214620
## SalePriceMM	2.406628e-03	0.0184235916		0.0087588669	11.3677068
## SalePriceCH	4.229308e-03	0.0025857973		0.0035977121	5.7424289
## PriceDiff	2.459433e-02	0.0505692331		0.0346948871	31.1857003
## Store7	9.799625e-04	0.0003121969		0.0007186679	0.8956352
## PctDiscMM	1.631394e-03	0.0026749205		0.0020242202	2.4066397
## PctDiscCH	-2.783524e-04	0.0022054421		0.0006976463	2.2417075
## ListPriceDiff	1.205162e-02	0.0068240479		0.0099998522	13.7020010
## STORE	8.605124e-03	0.0110309564		0.0095168506	9.1696040

Comments:

test error is 18.52% and the top3 most important predictors are:

"LoyalCH", "WeekofPurchase", "PriceDiff"

(n) Perform random forest on the training set with  $\sqrt{p}$  and  $p/3$  predictors respectively and report the prediction performance on the test set.

```
# with sqrt(p) ntry (sqrt(p) is the default for random forest classification)
# sqrt(p) is the default value of randomForest function:
set.seed(1)
rf.oj.train <- randomForest(Purchase~., data = train, importance = TRUE)

rf.pred.sqrt <- predict(rf.oj.train, newdata = test)
rf.table.sqrt <- table(rf.pred.sqrt, test$Purchase)
rf.error.sqrt <- 1-sum(diag(rf.table.sqrt))/sum(rf.table.sqrt)
print(paste("test error of the rf model with sqrt(p) mtry: ", round(rf.error.sqrt, 4)))
```



```
## [1] "test error of the rf model with sqrt(p) mtry: 0.1704"
```

```
# important predictors
print(rf.oj.train$importance)
```

##		CH	MM	MeanDecreaseAccuracy	MeanDecreaseGini
##	WeekofPurchase	0.0086483865	0.015784775	0.011523662	28.946406
##	StoreID	0.0096641735	0.031654548	0.018355206	21.499067
##	PriceCH	0.0045708309	0.003595271	0.004283876	5.639893
##	PriceMM	0.0019068217	0.007167926	0.004003358	6.272887
##	DiscCH	0.0003710221	0.004770353	0.002105822	4.262464
##	DiscMM	0.0042679863	0.009760813	0.006493291	5.355865
##	SpecialCH	0.0009488070	0.004243129	0.002279976	3.445623
##	SpecialMM	-0.0015845475	0.015044171	0.005037819	4.321306
##	LoyalCH	0.1131614535	0.225523308	0.157423714	155.678184
##	SalePriceMM	0.0078971280	0.018047829	0.011922103	11.362361
##	SalePriceCH	0.0048710638	0.004073971	0.004501026	7.106994
##	PriceDiff	0.0153258764	0.035382529	0.023292673	18.571671
##	Store7	0.0077114010	0.013011886	0.009804308	6.045269
##	PctDiscMM	0.0061440745	0.007567337	0.006711570	5.868678
##	PctDiscCH	0.0011998844	0.004153033	0.002336332	4.298745
##	ListPriceDiff	0.0088810210	0.013228920	0.010621118	10.781037
##	STORE	0.0070591358	0.021300943	0.012796562	12.273028

```
# with p/3 mtry
set.seed(1)
rf.oj.train2 <- randomForest(Purchase~., data = train, mtry=6, importance = TRUE)

rf.pred2<- predict(rf.oj.train2, newdata = test)
rf.table2 <- table(rf.pred2, test$Purchase)
rf.error2 <- 1-sum(diag(rf.table2))/sum(rf.table2)
print(paste("test error of the rf model with p/3 mtry: ", round(rf.error2, 4)))
```

```
## [1] "test error of the rf model with p/3 mtry: 0.1815"
```

```
# important predictors
print(rf.oj.train2$importance)
```

##		CH	MM	MeanDecreaseAccuracy	MeanDecreaseGini
##	WeekofPurchase	0.0110137766	0.015731401	0.012997396	32.928907
##	StoreID	0.0101491974	0.025989866	0.016299691	20.998314
##	PriceCH	0.0039803283	0.005347706	0.004561344	5.169630
##	PriceMM	0.0014647245	0.006856471	0.003588800	6.060659
##	DiscCH	0.0008635986	0.003304417	0.001826785	3.814710
##	DiscMM	0.0043140140	0.004177192	0.004290142	4.197832
##	SpecialCH	0.0014392566	0.003495084	0.002266185	4.080124
##	SpecialMM	-0.0011149369	0.007385660	0.002238215	4.069948
##	LoyalCH	0.1295831289	0.242840631	0.173901650	186.223445
##	SalePriceMM	0.0061637415	0.017952261	0.010753337	12.422317
##	SalePriceCH	0.0058933588	0.006389003	0.006061505	7.957959
##	PriceDiff	0.0177050757	0.037770237	0.025669676	21.000941
##	Store7	0.0062420816	0.009031380	0.007355507	5.453113

## PctDiscMM	0.0045882050	0.006137642	0.005233079	5.227907
## PctDiscCH	0.0002484615	0.004051293	0.001751392	3.804181
## ListPriceDiff	0.0073224034	0.010966501	0.008810964	11.517964
## STORE	0.0062344958	0.017638572	0.010781193	10.468506

Comments:

```
* for mtry = sqrt(p):
  test error is 17.04% and the top3 the most important predictors are:
  "LoyalCH", "WeekofPurchase", "StoreID"

* for mtry = p/3:
  test error is 18.15% and the top3 the most important predictors are:
  "LoyalCH", "WeekofPurchase", "PriceDiff"
```

(o) Compare the above models.

- 1) test error of Pruned Tree: 0.1704 the top predictor: "LoyalCH"
- 2) test error of the Boosting model with shrinkage parameter = 0.3: 0.1963 important predictors: "LoyalCH", "WeekofPurchase"
- 3) test error of the Bagging model: 0.1852" important predictors: "LoyalCH", "WeekofPurchase"
- 4) test error of the Random Forest with  $mtry = \sqrt{p}$ : 0.1704 important predictors are: "LoyalCH", "WeekofPurchase"
- 5) test error of the Random Forest with  $mtry = p/3$ : 0.1815 important predictors are: "LoyalCH", "WeekofPurchase"

Results: the Pruned Tree model and the Random Forest model with  $mtry=3/p$  produce the lowest test errors and their important predictors are the same